Concurrency Control in Distributed Database System

National Institute of Technology , Raipur

By

Ashish Verma(Update_10)

19111014

TERM PAPER SUBMISSION

feb 2022

Contents

1 ABSTRACT:

This paper reviews the coverage of concurrency control in Distributed networks. A distributed network becomes more popular, the need for improvement in distributed database management systems becomes even more important. The main challenges are identified as: (1)Preserving the ACID property atomicity, consistency, isolation, and durability property when concurrent transactions perform read and write operation; (2) providing recovery method when distributed data fail; (3)whatever method that is chosen they must provide feasible solutions with respect to performance. Keywords: Distributed Database, Distributed Design, Fragmentation, Replication, Allocation, Concurrency control, Transaction

1.1 INTRODUCTION:

requirement for secure, dependable, and open data in the present business climate, the requirement for conveyed information bases and customer/server applications is additionally expanding. A dispersed information base is a solitary coherent data set that is spread truly across PCs in different areas that are associated by information correspondence joins. A conveyed data set is a sort of virtual data set whose part parts are truly put away in various particular genuine data sets at various unmistakable areas. The clients at any area can get to information anyplace in the organization as though the information were completely put away at the client's own area. An appropriated information base administration framework is the product that deals with the Distributed Databases and gives an entrance component that makes this dispersion straightforward to the client. The goal of a circulated data set administration framework (DDBMS) is to control the administration of a disseminated data set (DDB) so that it appears to the client as a unified information base. This picture of a unified climate can be cultivated with the help of different sorts of straightforwardness like Location Transparency, Performance Transparency, Copy Transparency, Transaction Transparency, Transaction Transparency, Fragment Transparency, Schema Change Transparency,

and Local DBMS Transparency. Simultaneousness control is additionally a significant issue in data set frameworks. Simultaneousness control is the method involved with organizing simultaneous admittance to a data set in a multi-client information base administration framework (DBMS). There exist various strategies that give simultaneousness control. A portion of the strategies are Two-stage locking, Timestamping, Multi-adaptation timestamp, and so on.

2 MOTIVATION:

There are various business conditions that encourage the use of distributed databases: Data communications costs and reliability: If the data is geographically distributed and the applications are related to these data, it may be much more economical, in terms of communication costs, to partition the application and do the processing at each site. On the other hand, the cost of having smaller computing powers at each site is much less than the cost of having an equivalent power of a single mainframe.

Database recovery: The process of restoring data that has been lost, accidentally deleted, corrupted, or made inaccessible for any reason. In enterprise information technology (IT), data recovery typically refers to the restoration of data to a desktop, laptop, server, or external storage system from a backup.

Data sharing: the practice of making data used for scholarly research available to other investigators. Replication has a long history in science. The motto of The Royal Society is 'Nullius in verbal, translated "Take no man's word for it.

Distribution and autonomy of business units: Divisions, departments, and facilities in modern organizations are often geographically distributed, often across national boundaries. Often each unit has the authority to create

its own information systems, and often these units want local data over which they can have control. Business mergers and acquisitions often create this environment.

Improved Performance: Performance improvement, by its nature, is iterative. For this reason, removing the first bottleneck might not lead to performance improvement immediately, because another bottleneck might be revealed. Also, in some cases, if serialization points move to a more inefficient sharing mechanism,then performance could degrade. With experience, and by following a rigorous method of bottleneck elimination, applications can be debugged and made scalable.

Increased reliability and availability: When a centralized system fails, the database is unavailable to all users. In contrast to centralized systems, the distributed systems will continue to function at some reduced level, however, even when a component fails. Faster response: Well-suited for a loosely defined data structure that may evolve over time. In-memory database management system (IMDBMS) - provides faster response times and better performance.

3   DISTRIBUTED DATABASE DESIGN:

Distributed databases can be homogenous or heterogeneous. In a homogeneous distributed database system, all the physical locations have the same underlying hardware and run the same operating systems and database applications. In a heterogeneous distributed database, the hardware, operating systems, or database applications may be different at each of the locations

Data Fragmentation: In Distributed Databases, we need to define the logical unit of Database Distribution and allocation. The database may be broken up into logical units called fragments which will be stored at different sites. The simplest logical units are the tables themselves.

3.1 Three Types of Data Fragmentation are:

- Horizontal fragmentation:refers to the division of a relation into subsets (fragments) of tuples (rows).

  Example: Say we have this relation

customer_id	Name	Area	Payment Type	sex
1	Kashish	Chhattisgarh	debit card	female
2	Ragini	Goa	Credit card	female

Horizontal Fragmentation are subsets of tuples (rows)

Fragment 1

customer_id	Name	Area	Payment Type	sex
1	Kashish	Chhattisgarh	debit card	female
2	Ragini	Goa	Credit card	female

Fragment 2

customer_id	Name	Area	Payment Type	sex
3	Manan	Punjab	debit card	male

- Vertical fragmentation:refers to the division of a relation into subsets (fragments) of tuples (rows).Vertical Fragmentation of a relation R produces fragments R1, .., Rn, each of.

  Example: Say we have this relation

Vertical Fragmentation are subsets of tuples (rows)

Fragment 1

customer_id	Name	Area	sex
1	Kashish	Chhattisgarh	female
2	Ragini	Goa	female
3	Manan	Punjab	male

Fragment 2

customer_id	Payment type
1	debit card
2	credit card
3	debit card

- Hybrid fragmentation: Hybrid Fragmentation comprises the combination of characteristics of both Horizontal and Vertical Fragmentation. Each fragment can be specified by a SELECT-PROJECT combination of operations. In this case, the original table can be reconstructed by applying union and natural join operations in the appropriate order

Data Replication: the frequent electronic copying of data from a database in one computer or server to a database in another so that all users share the same level of information. The result is a distributed database in which users can access data relevant to their tasks without interfering with the work of others.

4   FUNDAMENTALS OF TRANSACTION MANAGEMENT AND CONCURRENCY CONTROL:

4.1   Transaction:

is a sequence of operations performed as a single logical unit of work. A logical unit of work must exhibit four properties, called the atomicity, consistency, isolation, and durability (ACID) properties, to qualify as a transaction.

4.2   Properties of Transaction:

A Transaction has four properties that lead to the consistency and reliability of a distributed database. These are Atomicity, Consistency, Isolation, and Durability.

Atomicity: An atomic transaction is an indivisible and irreducible series of database operations such that either all occur, or nothing occurs.

Consistency: database systems refer to the requirement that any given database transaction must change affected data only in allowed ways. Any data written to the database must be valid according to all defined rules, including constraints, cascades, triggers, and any combination thereof.

Isolation: in the context of databases, specifies when and how the changes implemented in an operation become visible to other parallel operations. Transaction isolation is an important part of any transactional system. It deals with consistency and completeness of data retrieved by queries unaffecting a user data by other user actions. A database acquires locks on data to maintain a high level of isolation.

Durability: databases are the property that ensures transactions are saved permanently and do not accidentally disappear or get erased, even during a database crash. This is usually achieved by saving all transactions to a nonvolatile storage medium.

4.3 Concurrency Control:

A database management system (DBMS) concept that is used to address conflict with the simultaneous accessing or altering of data that can occur with a multiuser system.

5 DISTRIBULATED CONCURRENCY CONTROL ALGORITHMS:

In this paper, we consider some of the distributed concurrency control algorithms. We summarize the salient aspects of these four algorithms in this section. In order to do this, we must first explain the structure that we have

assumed for distributed transactions. Before discussing the algorithms, we need to get an idea about the distributed transactions. Distributed Transaction: A distributed transaction is a transaction that runs in multiple processes, usually on several machines. Each process works for the transaction. Distributed transaction processing systems are designed to facilitate transactions that span heterogeneous, transaction-aware resource managers in a distributed environment. The execution of a distributed transaction requires coordination between a global transaction management system and all the local resource managers of all the involved systems. The resource manager and transaction processing monitor are the two primary elements of any distributed transactional system.

Distributed transactions, like local transactions, must observe the ACID properties. However, maintenance of these properties is very complicated for distributed transactions because a failure can occur in any process. If such a failure occurs, each process must undo any work that has already been done on behalf of the transaction.

A distributed transaction processing system maintains the ACID properties in distributed transactions by using two features:

Recoverable processes: Recoverable processes log their actions and therefore can restore earlier states if a failure occurs.

A commit protocol: A commit protocol allows multiple processes to coordinate the committing or aborting of a transaction. The most common commit protocol is the two-phase commit protocol.

Distributed Two-Phase Locking (2PL): Two-phase locking (2PL) is a concurrency control method that guarantees serializability. It is also the name of the resulting set of database transaction schedules (histories). The protocol utilizes locks, applied by a transaction to data, which may block (interpreted as signals to stop) other transactions from accessing the same data during the transaction's life.

The 2PL Protocol oversees locks by determining when transactions can acquire and release locks. The 2PL protocol forces each transaction to make a lock or unlock request in two steps:

Growing Phase(Expanding phase): locks are acquired and no locks are released.

Shrinking Phase: locks are released and no locks are acquired.

The transaction first enters into the Growing Phase, makes requests for required locks, then gets into the Shrinking phase where it releases all locks and cannot make any more requests. Transactions in 2PL Protocol should get all needed locks before getting into the unlock phase. While the 2PL protocol guarantees serializability, it does not ensure that deadlocks do not happen. So deadlock is a possibility in this algorithm, Local deadlocks are checked for any time a transaction blocks, and are resolved when necessary by restarting the transaction with the most recent initial startup time among those involved in the deadlock cycle. Global deadlock detection is handled by a "Snoop" process, which periodically requests waits-for information from all sites and then checks for and resolves any global deadlocks.

6 Wound-Wait (WW):

The second algorithm is It is a preemptive technique for deadlock prevention. It is a counterpart to the wait-die scheme. When Transaction Ti requests a data item currently held by Tj, Ti is allowed to wait only if it has a timestamp larger than that of Tj, otherwise Tj is rolled back (Tj is wounded by Ti). In this scheme, if a transaction requests to lock a resource (data item), which is already held with conflicting lock by some another transaction, one of the two

possibilities may occur.

Example: wound wait algorithm

" "	T1 is allowed to
t(T1) > t(T2)	wait
t(T1) < t(T2)	Abort and rollrd back

t(T1) > t(T2) -: If requesting transaction [t(T1)] is younger than the transaction [t(T2)] that has holds lock on requested data item then requesting transaction [t(T1)] has to wait.

t(T1) < t(T2) -: If requesting transaction [t(T1)] is older than the transaction [t(T2)] that has holds lock on requested data item then requesting transaction [t(T1)] has to abort or rollback.

Suppose that Transactions T22, T23, T24 have time-stamps 5, 10 and 15 respectively . If T22requests a data item held by T23, then data item will be preempted from T23 and T23 will be rolled back. If T24 requests a data item held by T23, then T24 will wait.

Basic Timestamp Ordering (BTO): We assume that the Transaction Manager (TM) attaches an appropriate timestamp to all read and write operations. In the BTO, the scheduler at each Data Manager (DM), keeps track of the largest timestamp of any read and write operation processed thus far for each data object. These timestamps may be denoted by R-ts(object) and W-ts(object), respectively. Let us also make the following notations:

read(x,TS) –> Read request with timestamp TS on a data object x.

write(x,v,TS) –> Write request with timestamp TS on a data object x. v is the

value to be assigned to x.

A read request is handled in the following manner:

If TS < W-ts(x) then reject read request and abort corresponding transaction

Else

execute transaction

11

Set R-ts(x) to maxR-ts(x), TS

A write request is handled in the following manner:

If TS < R-ts(x) or TS < W-ts(x) then reject write request

else

execute transaction

Set W-ts(x) to TS

If a transaction is aborted, it is restarted with a new timestamp. This can result in a cyclic restart where a transaction can repeatedly restart and abort without ever completing. Another disadvantage is that it has storage overhead for maintaining timestamps (two timestamps must be kept for every data object).

7   Distributed Optimistic(OPT):

The fourth algorithm is Locking bad because :

- Overhead not present in sequential case. Even read-only transactions need to use locks.

- Lack of general-purpose deadlock free locking protocols.

- Large parts of DB on secondary storage - concurrency significantly lowered when it is necessary to leave congested node while accessing disk

- to allow transaction to abort, need to hold on to locks until EOT

- locking may only be necessary in worst case.

1) Optimistic: reading a value/pointer can never cause loss of integrity; reads are completely unrestricted, although returning a value from a query is considered equivalent to a write.

2) writes are severly restricted. Have read phase, validation phase, then write phase. During read, all writes are on local copies. Validation ensures consistency across active transactions. Write phase is when local copies get written

to global DB.

3) Concurrency control procedures maintain sets of objects accessed by transaction. maintain read and write set for transaction.

4) Serial equivalence: if have individual transactions that are executed concurrently, have serial equivalence if there exists some serial sequence of the transactions that produce the final data structure. serial equivalence is sufficient (but not necessary) for integrity of DB.

5) Validation of serial equivalence: assign each transaction a transaction number (TN), and explicitly force serializability. Need one of these three conditions.

6) Ti completes write phase before Tj starts read phase.

7) Write set of Ti does not intersect read set of Tj, and Ti completes its write phase before Tj starts its write phase.

8) Write set of Ti does not intersect read set or write set of Tj, and Ti completes its read phase before Tj completes its read phase. (1) says Ti completes before Tj starts. (2) states writes of Ti don't affect read phase of Tj, and Ti finishes writing before Tj starts writing, hence doesn't overwrite Tj. (3) similar to (2), but does not require that Ti finish writing before Tj starts writing.

8   Main problems in using Concurrency

The problems which arise while using concurrency are as follows:

Updates will be lost   One transaction does some changes and another transaction deletes that change. One transaction nullifies the updates of another transaction.

Uncommitted Dependency or dirty read problem _ On variable has updated in one transaction, at the same time another transaction has started and deleted the value of the variable there the variable is not getting updated or committed that has been done on the first transaction this gives us false values or the previous values of the variables this is a major problem.

Inconsistent retrievals _ One transaction is updating multiple different variables, another transaction is in a process to update those variables, and the

problem occurs is inconsistency of the same variable in different instances.

9   Concurrency control techniques

The concurrency control techniques are as follows:-

Locking   Lock guaranties exclusive use of data items to a current transaction. It first accesses the data items by acquiring a lock, after completion of the transaction it releases the lock.

Types of Locks:

The types of locks are as follows:

Shared Lock [Transaction can read only the data item values]
Exclusive Lock [Used for both read and write data item values]

Time Stamping   Time stamp is a unique identifier created by DBMS that indicates relative starting time of a transaction. Whatever transaction we are doing it stores the starting time of the transaction and denotes a specific time.

This can be generated using a system clock or logical counter. This can be started whenever a transaction is started. Here, the logical counter is incremented after a new timestamp has been assigned.

Optimistic   It is based on the assumption that conflict is rare and it is more efficient to allow transactions to proceed without imposing delays to ensure serializability.

Assign transactions numbers when transaction enters validation, to prevent fast transactions from blocking behind slow but earlier starting transactions. 1) Serial validation: assign transaction nmber, validate, and do write

phase all in critical section. Fine if one CPU and write phase can take place in main memory. Otherwise too inefficient - not enough concurrency. If multiple CPUs, be more clever about critical sections. 2) Parallel validation: maintain set of active transactions (in read but not yet completed write phase); do validation against all transactions in active set. To prevent an aborted transaction from causing further aborts, do multistage validation. 3) Analysis for large B-trees shows that optimistic concurrency control is a win -very rare that one insertion would cause another concurrent insertion to restart. (Assumes random access to B-tree).